

DMF for Crystal Structure Prediction

One-shot generative model matches diffusion CSP at 40× lower compute (for half of MP-20)

PoC adapting Drifting Models with Friction to MP-20.

Setup & TL;DR

Task: given composition (atom types, N), predict crystal structure (lattice + frac coords).

Benchmark: MP-20 test, pymatgen StructureMatcher (ltol=0.3, stol=0.5, angle=10°).

Key metric — match@K:

- For each test composition, generate **K** candidate structures.
- match@K = % of test compositions where **at least one** of the K samples matches the ground-truth structure within tolerance.
- DiffCSP-CSP paper standard: K=20. Compute cost = K × NFE-per-sample.

Baseline: DiffCSP — diffusion, 1000 NFE/sample, ~81% match@20 aggregate.

DMF idea: one-shot generator with kernel-mean-shift "drift" loss → **1 NFE/sample**, so we can afford much larger K at the same compute budget.

Headline: DMF + a **repulsion term** in the drift loss matches DiffCSP for N=2-8 atoms-per-cell at **40× lower compute**.
Hard cliff at $N \geq 9$.

DMF + repulsion in V

Per training step:

1. Sample noise $z \sim N(0, I)$, one forward through CSPNet $\rightarrow (\hat{L}, \hat{F})$
2. Compute drift V from a kernel-mean-shift over the target batch
3. Add explicit **anti-mode-seeking term**:

$$V_{\text{total}} = V_{\text{attract}}(\text{gen} \rightarrow \text{target_batch}) - \lambda \cdot V_{\text{repel}}(\text{gen} \rightarrow \text{other_gen})$$

4. Loss = MSE between (gen, gen + V_{total})

Single-line code change in `v_drift.py`. No iterative refinement. No 1000 reverse-diffusion steps. Just one shot per sample.


Without repulsion: mode-collapse $\rightarrow$ match@20 = 10.6%. With repulsion $\lambda=1.0$: **26.6%** (+150%).

λ_{rep} sweep: optimum at 1.0

λ_{rep}	match@20 (limit=500)	validity
0.0 (no repulsion)	10.6%	0.253
0.5	16.8%	0.408
1.0	26.6%	0.349
1.5	24.6%	0.397 (balanced)
2.0	20.8%	(lattices degenerate)


Convex curve. $\lambda=1.0$ — **operational default** for match-rate. $\lambda=1.5$ — balanced default (high validity + near-peak match).

Compute-parity head-to-head



Same compute budget (5000 NFE/composition): **DMF wins for N=4-8** by 4-25 pp over DiffCSP. N=5: 100% vs 75% (+25 pp). N=6: 94% vs 83% (+11 pp).

Per-N saturation with K



DMF's per-N match-rate **crosses DiffCSP K=20 baseline** (dotted) as K grows.

N=4: by K=100. N=6: by K=500. N=8: by K=500 (tied with DiffCSP). N=10: barely (7.7%). N=12: zero — hard architectural cliff.

Final per-N table (K=500)

N	n	DMF K=500	DiffCSP K=20	DMF wins?
2	3	100%	86%	✓ +14
3-5	36	100%	95-100%	✓
6	18	100%	84%	✓ +16
7	10	90%	79%	✓ +11
8	25	84%	84%	TIED
9	5	20%	82%	✗
10	26	7.7%	90%	✗
12+	many	0%	70-90%	✗ hard cliff

DMF wins or ties for N=2-8 (≈half of MP-20 by count) at **40× lower compute** (500 vs 20,000 NFE/comp).

DMF is intrinsically divergent

For each composition, count distinct structural clusters among 20 samples:

Method	mean clusters	%mode-collapsed	% ≥ 3 distinct
DiffCSP	5.88	24%	61%
DMF + repulsion	17.89	0%	100%

DMF gives **3× more distinct candidates per composition**, zero collapse.

DiffCSP = convergent (1000 reverse steps cluster to mode). DMF = divergent (one-shot $z \rightarrow$ diverse outputs). **Different paradigms**, not just different speed points.

→ DMF strictly better for **screening pipelines** (cheaper, more diverse candidates for downstream filters like CHGNet/DFT).

Why it works / why $N \geq 9$ fails

Works for small- N because:

- Repulsion forces samples to spread $\rightarrow$ high diversity per composition
- 1-NFE per sample $\times$ 220 $\times$ speedup vs DiffCSP $\times$ 3 $\times$ diversity = many cheap candidates
- For $N \leq 8$, multi-sample coverage covers GT structure within tolerance

Fails for $N \geq 9$ because:

- One-shot architecture cannot organize 9+ atoms into coherent periodic cell
- Multi-atom coordination requires iterative refinement (diffusion's strength)
- $N=12+$: literally zero matches at $K=500$ — model **cannot produce** these structures

Architectural limit, not a tuning issue. Open problem.

Failed approaches (8 of 9 ablations net-negative)

Before the loss-side fix, these all hurt match-rate:

- Lattice prior tuning, V temperature sweep
- K-step iterative training (mean-seeking)
- Per-atom z (destroys inter-atom coherence)
- FiLM composition conditioning (tighter RMSE, worse match-rate)
- Memory bank, no-friction training

The single loss-side intervention (repulsion) broke through where 8 architecture/hyperparameter knobs failed.

Lesson: **architecture is downstream of loss.** Don't tune the model when the objective is wrong.

Summary & practical use

	DiffCSP	DMF+repulsion
Inference per sample	1000 NFE	1 NFE
Inference 10k crystals	~30 min	8.3 sec
Match@20 aggregate (limit=500)	84%	27% (K=20), 45% (K=500)
Match@20 per-N (N=2-8)	79-100%	84-100% (K=500, ties or wins)
Distinct clusters per comp	5.88	17.89
Works for $N \geq 9$	✓	✗

Recommended use:

- **Small-cell CSP / screening pipelines** → DMF (40× cheaper, 3× more diverse)
- **Large-cell prediction ($N \geq 9$)** → DiffCSP
- **Mixed workloads** → hybrid pipeline

Status: 25 sub-tasks completed, full reports + 3 paper-ready figures in `dmf_poc/` . Ready for paper draft.